

Totem Interactive

ML Developer Assignment

Dynamic LLM Router & RAG System

Design a **plug-and-play** solution that **routes** user requests to different LLM providers based on **request type** (e.g., “technical,” “creative,” “general”). In parallel, it **asynchronously** queries a **low-fidelity model** to ensure there’s always a fallback response if the primary model fails. The system must also **log all interactions** and **dynamically update** each LLM’s priority based on real-world performance. Additionally, allow a mechanism to **search for the “best” LLaMA model** for a given use case and append it to the active configuration.

Core Objectives

1. Request Categorization & LLM Selection

- Classify each user query (e.g., “technical,” “creative,” “general”) using basic NLP or a rule-based approach.
- Map each category to a **preferred** LLM (e.g., a creative LLM for “creative,” a more factual LLM for “technical,” etc.).

2. Asynchronous Fallback Model

- In parallel with the “preferred” model call, **fire a request** to a **lower-fidelity** (or cheaper) model.
- If the preferred LLM times out or fails, **immediately return** the fallback model’s response.
- If the preferred LLM’s response arrives successfully after the fallback, you can update the conversation or just log it for reference.

3. Performance-Based Dynamic Priority

- **Log** all responses (success, time taken, cost, etc.).
- Adjust the **priority** of each LLM **per category** based on observed performance (e.g., if an LLM consistently fails for “technical,” move it down in the priority list).

4. RAG Application (Low-Level) & Model Extension

Totem Interactive

ML Developer Assignment

- Treat the fallback model as a **Retrieval-Augmented Generation** (RAG) style or a simple knowledge-based model to ensure minimal coverage.
- Implement a “**search for the best LLaMA model**” step that can discover new LLaMA variants (or other models) for a category, then **append** them to the config if they seem promising.
- This search could be a mock or partial implementation that simulates calling an API or scanning a repository for the “best” LLaMA model.

5. Logging & Dynamic Config Updates

- Store logs with details of each request:
 - **Category** assigned
 - **Preferred LLM** used
 - **Fallback LLM** used (if any)
 - **Response time, Success/Fail, and Cost** (if available).
- Based on logs, **dynamically** reassign the best model for each category (e.g., “technical” => Model B if Model A is failing).

6. Demo & Real-World Testing

- Prepare a small **demo UI** or CLI to show how the routing works for different query types.
- The final step involves **replacing** a mock LLaMA setup with a **real LLaMA API key** (provided later in-person) and testing **debugging** or **live performance** in your environment.

Mandatory Features

1. Categorization Logic

- At least a **simple** classification approach to label queries by topic or style.

Totem Interactive

ML Developer Assignment

2. Asynchronous Requests

- Fire calls to multiple models simultaneously or in parallel threads/processes.

3. Fallback Mechanism

- Immediately serve the fallback model's output if the primary fails or lags beyond a set timeout.

4. Dynamic Priority Updates

- Re-rank LLM providers over time using logs.
-

Deliverables

1. Source Code (Microservice or Library)

- A **well-structured** repository (e.g., GitHub) with a **clear README**.
- Include **installation** and **run** steps (Docker or local environment).

2. Configuration & Categorization Files

- A config for LLM providers (endpoints, cost, priority per category).
- Basic rules or an ML model for categorizing requests (technical, creative, etc.).

3. Logging & Dynamic Updates

- Show logs for each request in a structured format (JSON, database, or console).
- Demonstrate how the system reorders LLM priorities over time based on logs.

4. Search & Append Mechanism

- A **mock or minimal** function that “searches” for a new LLaMA variant and appends it to the config if it passes some criteria.

Totem Interactive

ML Developer Assignment

- This can be partial, but show the logic of how it would integrate.

5. Demo UI

- A **simple interface** (web page or CLI) where you can enter a prompt, see which model(s) got called, the fallback model if used, and the final response.
-

Evaluation Criteria

1. Technical Implementation (40%)

- Proper asynchronous calls, fallback logic, dynamic priority re-ranking.
- Quality of request categorization approach and overall code structure.

2. Resilience & Logging (30%)

- Robust error/timeout handling for the main vs. fallback model calls.
- Thorough logs that enable dynamic updates to the model priority.

3. Search & Append Feature (20%)

- Creativity and clarity in simulating or partially implementing a “search” for new LLaMA variants.
- Proper integration into the main config so the new model can be used if it's better for a certain category.

4. Demo & Usability (10%)

- A straightforward UI/CLI to show how the system reacts to different query types.
 - Clarity in how developers can integrate or extend this system.
-

Submission Guidelines

Totem Interactive

ML Developer Assignment

1. GitHub Repository

- Name it, for example, `Dynamic-LLM-Router-{YourName}`.
- Include separate folders for core code (`/src` or `/app`), config files, logs (if needed), and your demo UI.

2. README.md

- Detail setup instructions, dependencies, usage examples.
- Explain assumptions, limitations, and any special instructions for running asynchronous tasks.

3. Commit & Push

- Ensure version control with clear commit messages.
- Provide instructions on how to run the final service or test script.

4. Live Integration Step (Optional but Encouraged)

- After the baseline is done, be prepared to **replace** your mock LLaMA integration with a **real** LLaMA API key (provided in person) and show debugging or performance testing.

Submit your GitHub repo URL and any additional notes (e.g., how to replicate the search & append feature if it requires specific external services). We'll review your code and potentially test it with our own keys to see how well it adapts in real time.

Objective:

Create an **LLM router** that intelligently **assigns requests** by category, uses a **fallback** model asynchronously, **logs everything**, and **updates priorities** dynamically based on performance. Demonstrate a method to **search and add** new LLaMA models, ensuring your system remains **flexible and scalable** for evolving AI needs.

Notes

Totem Interactive

ML Developer Assignment

- The implementation approach is open-ended; focus on fulfilling the core requirements.
- Prioritize the quality of analysis logic over interface aesthetics.
- Feel free to include additional features if they enhance functionality or user experience.

Submission

Submit the following files to career@toteminteractive.in / Totemistaken@gmail.com / Shoeb@toteminteractive.in / arjungujar@toteminteractive.in:

1. Create a GitHub repository for your assignment solution
 - Provide a descriptive name (e.g., "{Name}-assignment-solution")
 - Create a folder for your code (e.g., src or code)
 - Add solution files (e.g., .swift, .py, .js in the coding language of your choice)
 - Use Git for version control and commit changes
2. Create a README.md file in the root directory
 - Explain implementation details and approach
 - Provide usage instructions
 - Mention assumptions, limitations, and any identified edge cases
 - List dependencies and additional notes
3. Commit and push all changes to the remote repository
4. Submit the GitHub repository URL to the designated contact person